

Data Cleaning:

1. Remove unnecessary columns:

```
scala> import org.apache.spark.sql.SparkSession
import org.apache.spark.sql.SparkSession

scala> val spark = SparkSession.builder().appName("Remove Columns").getOrCreate()
```

- Read in the original csv file. We can see it has over 37 columns.

```
scala> val df = spark.read.option("header","true").option("inferSchema","true").csv("game_data.csv.csupload")
df: org.apache.spark.sql.DataFrame = [AppID: int, Name: string ... 37 more fields]
```

- If we print them, there are all the columns.

```
scala> df.printSchema()
root
 |-- AppID: integer (nullable = true)
 |-- Name: string (nullable = true)
 |-- Release date: string (nullable = true)
 |-- Estimated owners: string (nullable = true)
 |-- Peak CCU: integer (nullable = true)
 |-- Required age: integer (nullable = true)
 |-- Price: double (nullable = true)
 |-- DLC count: integer (nullable = true)
 |-- About the game: string (nullable = true)
 |-- Supported languages: string (nullable = true)
 |-- Full audio languages: string (nullable = true)
 |-- Reviews: string (nullable = true)
 |-- Header image: string (nullable = true)
 |-- Website: string (nullable = true)
 |-- Support url: string (nullable = true)
 |-- Support email: string (nullable = true)
 |-- Windows: boolean (nullable = true)
 |-- Mac: boolean (nullable = true)
 |-- Linux: boolean (nullable = true)
 |-- Metacritic score: integer (nullable = true)
 |-- Metacritic url: string (nullable = true)
 |-- User score: integer (nullable = true)
 |-- Positive: integer (nullable = true)
 |-- Negative: integer (nullable = true)
 |-- Score rank: double (nullable = true)
 |-- Achievements: integer (nullable = true)
 |-- Recommendations: integer (nullable = true)
 |-- Notes: string (nullable = true)
 |-- Average playtime forever: integer (nullable = true)
 |-- Average playtime two weeks: integer (nullable = true)
 |-- Median playtime forever: integer (nullable = true)
 |-- Median playtime two weeks: integer (nullable = true)
 |-- Developers: string (nullable = true)
 |-- Publishers: string (nullable = true)
 |-- Categories: string (nullable = true)
 |-- Genres: string (nullable = true)
 |-- Tags: string (nullable = true)
 |-- Screenshots: string (nullable = true)
 |-- Movies: string (nullable = true)
```

- Then I delete some columns in the csv file

```
scala> val columnsToKeep = df.columns.filterNot(column => List("Reviews", "Website", "Support url", "Support email", "Metacritic url", "Score rank", "Notes", "Tags", "Screenshots", "Movies").contains(column))
columnsToKeep: Array[String] = Array(AppID, Name, Release date, Estimated owners, Peak CCU, Required age, Price, DLC count, About the game, Supported languages, Full audio languages, Header image, Windows, Mac, Linux, Metacritic score, User score, Positive, Negative, Achievements, Recommendations, Average playtime forever, Average playtime two weeks, Median playtime forever, Median playtime two weeks, Developers, Publishers, Categories, Genres)
```

- We can see there are 27 columns right now.

```
scala> val refinedDF = df.select(columnsToKeep.map(df.col):_*)
refinedDF: org.apache.spark.sql.DataFrame = [AppID: int, Name: string ... 27 more fields]

scala> refinedDF.printSchema()
root
 |-- AppID: integer (nullable = true)
 |-- Name: string (nullable = true)
 |-- Release date: string (nullable = true)
 |-- Estimated owners: string (nullable = true)
 |-- Peak CCU: integer (nullable = true)
 |-- Required age: integer (nullable = true)
 |-- Price: double (nullable = true)
 |-- DLC count: integer (nullable = true)
 |-- About the game: string (nullable = true)
 |-- Supported languages: string (nullable = true)
 |-- Full audio languages: string (nullable = true)
 |-- Header image: string (nullable = true)
 |-- Windows: boolean (nullable = true)
 |-- Mac: boolean (nullable = true)
 |-- Linux: boolean (nullable = true)
 |-- Metacritic score: integer (nullable = true)
 |-- User score: integer (nullable = true)
 |-- Positive: integer (nullable = true)
 |-- Negative: integer (nullable = true)
 |-- Achievements: integer (nullable = true)
 |-- Recommendations: integer (nullable = true)
 |-- Average playtime forever: integer (nullable = true)
 |-- Average playtime two weeks: integer (nullable = true)
 |-- Median playtime forever: integer (nullable = true)
 |-- Median playtime two weeks: integer (nullable = true)
 |-- Developers: string (nullable = true)
 |-- Publishers: string (nullable = true)
```

- Then I download the file back to HDFS

```
refinedDF.write.option("header", "true").csv("cleaned_game_data.csv")
```

2. Remove null elements in the csv file:

- Similar to the first one, I read the csv file from the first one. Then I do

```
scala> val df = spark.read.option("header", "true").option("inferSchema", "true").csv("cleaned_game_data.csv")
df: org.apache.spark.sql.DataFrame = [AppID: int, Name: string ... 27 more fields]
```

- Then I just use this code to drop rows that have null values.

```
scala> val nonNullDF = df.na.drop()
nonNullDF: org.apache.spark.sql.DataFrame = [AppID: int, Name: string ... 27 more fields]
```

- Then I write the file back to HDFS, and I get a cleaned dataset, called:
cleaned_game_data_2.csv

```
scala> nonNullDF.write.option("header", "true").csv("cleaned_game_data_2.csv")
```

Data Analysis

1. Correlation between features

- We can see there is quite a high correlation between the “Positive” view and “Peak CCU”. There is little correlation between price and user score, price and Negative reviews.

```
scala> val processedData = data.withColumn("Peak CCU", col("Peak CCU").cast("int")).withColumn("Positive",  
ive").cast("int")).withColumn("Price", col("Price").cast("float"))  
processedData: org.apache.spark.sql.DataFrame = [AppID: string, Name: string ... 27 more fields]  
  
scala> val correlationData = processedData.stat.corr("Positive", "Peak CCU")  
correlationData: Double = 0.7101076163637777  
  
scala> println(s"Correlation between positive reviews and peak concurrent users: $correlationData")  
Correlation between positive reviews and peak concurrent users: 0.7101076163637777
```

```
a> val corrPriceUserScore = processedData.stat.corr("Price", "User score")  
PriceUserScore: Double = -0.0038111685236258946  
  
a> val corrPriceNegative = processedData.stat.corr("Price", "Negative")  
PriceNegative: Double = 0.030703322812391368  
  
a> val corrPositiveNegative = processedData.stat.corr("Positive", "Negative")  
PositiveNegative: Double = 0.7294850839832202
```

- Also, there is little correlation between price and the time a player spends on a game. Also, longer a player spends on a game doesn't mean a positive review.

```
scala> val corrPlaytimePrice = processedData.stat.corr("Average playtime forever", "Price")  
corrPlaytimePrice: Double = 0.0829825953304325  
  
scala> val corrPlaytimePositive = processedData.stat.corr("Average playtime forever", "Positive")  
corrPlaytimePositive: Double = 0.08088311491473374  
  
scala> []
```

- I checked the “User score” column, and saw each values' counts. Then I found that most data is zero, so this feature is also not very useful.

```
scala> val spark = SparkSession.builder().appName("User score Frequency Analysis").master("local[*]").getOrCreate()  
24/04/27 16:38:46 WARN org.apache.spark.sql.SparkSession$Builder: Using an existing SparkSession; some spark core co  
spark: org.apache.spark.sql.SparkSession = org.apache.spark.sql.SparkSession@128e7fc4  
  
scala> val data = spark.read.option("header", "true").option("inferSchema", "true").csv("cleaned_game_data_2.csv")  
data: org.apache.spark.sql.DataFrame = [AppID: int, Name: string ... 27 more fields]  
  
scala> val processedData = data.withColumn("User score", col("User score").cast("int"))  
processedData: org.apache.spark.sql.DataFrame = [AppID: int, Name: string ... 27 more fields]  
  
scala> val scoreFrequency = processedData.groupBy("User score").count().orderBy("User score")  
scoreFrequency: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [User score: int, count: bigint]  
  
scala> scoreFrequency.show()  
+-----+-----+  
|User score|count|  
+-----+-----+  
|         0|11664|  
|         51|    1|  
|         57|    1|  
|         59|    1|  
|         63|    1|  
|         68|    1|  
|         69|    1|  
|         77|    1|  
|         80|    2|  
|        100|    2|  
+-----+-----+
```

- I created a new column called “positive_ratio”, which calculate the ratio of positive votes.

```
scala> val processedData = data.withColumn("Positive",col("Positive").cast("int")).withColumn("Negative",col("Negative").cast("int"))
processedData: org.apache.spark.sql.DataFrame = [AppID: int, Name: string ... 27 more fields]

scala> val withPositiveRatio = processedData.withColumn("positive_ratio",col("Positive")/(col("Positive")+col("Negative")+0.0001))
withPositiveRatio: org.apache.spark.sql.DataFrame = [AppID: int, Name: string ... 28 more fields]

scala> withPositiveRatio.select("positive_ratio").show(10)
+-----+
| positive_ratio|
+-----+
| 0.999990000099999|
|0.7851233180799024|
|0.7049168771854472|
| 0.999987500156248|
|0.8070737354358754|
|0.9073632049018516|
|0.9433953364194939|
|0.999993333777775|
| 0.249996875039062|
|0.7016123374094053|
+-----+
only showing top 10 rows
```

-